

# 02285 AI and MAS, F23

## Conflict based search

Today's subjects:

- Multi-agent pathfinding (MAPF)
- Baseline algorithm: Multibody A\*
- Explore three (and a half) other algorithms
- In particular, conflict based search

# MAPF problem

## Multi-agent pathfinding (MAPF) problem


Given a problem instance

- a graph  $G = (V, E)$ , where vertices  $V$  of the graph are possible locations for the agents, and the edges  $E$  are the possible transitions between locations;
- a set  $a_1, a_2, \dots, a_k$  of  $k$  agents;
- a set  $v_{1,0}, v_{2,0}, \dots, v_{k,0}$  of start locations for each agent;
- a set  $g_1, g_2, \dots, g_k$  of goal locations for each agent;

A solution is a plan  $\pi_i = (v_{i,0}, v_{i,1}, \dots, g_i)$  for each agent  $a_i$ , i.e., a sequence of vertices that lead from the start to the goal.

# MAPF problem

- Simpler than the hospital domain
  - No boxes
  - All agents have one goal
  - But considers general graphs, not only grids
- Today we will look at different algorithms for solving this problem



## Solution 1: Multibody A\*

- A\* graph search algorithm in the space of joint actions, i.e., multibody A\* (MBA\*)
- Similar to what you did in the warmup assignment

## Solution 1: Multibody A\*

- A\* graph search algorithm in the space of joint actions, i.e., multibody A\* (MBA\*)
- Similar to what you did in the warmup assignment
- Heuristic:

$$h(s) = \sum_{i=1}^k \text{dist}(v_{i,s}, g_i)$$

where

- $v_{i,s}$  is the vertex where agent  $a_i$  is located in state  $s$ , and
- $\text{dist} : V \mapsto V$  is the shortest distance from any vertex to any other vertex (can be pre-computed, e.g., using the Floyd-Warshall algorithm in  $O(|V|^3)$  time)

## Solution 1: Multibody A\*

- A\* graph search algorithm in the space of joint actions, i.e., multibody A\* (MBA\*)
- Similar to what you did in the warmup assignment
- Heuristic:

$$h(s) = \sum_{i=1}^k \text{dist}(v_{i,s}, g_i)$$

where

- $v_{i,s}$  is the vertex where agent  $a_i$  is located in state  $s$ , and
- $\text{dist} : V \mapsto V$  is the shortest distance from any vertex to any other vertex (can be pre-computed, e.g., using the Floyd-Warshall algorithm in  $O(|V|^3)$  time)
- MBA\* is complete (finds a solution if there is one)

## Solution 1: Multibody A\*

- A\* graph search algorithm in the space of joint actions, i.e., multibody A\* (MBA\*)
- Similar to what you did in the warmup assignment
- Heuristic:

$$h(s) = \sum_{i=1}^k \text{dist}(v_{i,s}, g_i)$$

where

- $v_{i,s}$  is the vertex where agent  $a_i$  is located in state  $s$ , and
- $\text{dist} : V \mapsto V$  is the shortest distance from any vertex to any other vertex (can be pre-computed, e.g., using the Floyd-Warshall algorithm in  $O(|V|^3)$  time)
- MBA\* is complete (finds a solution if there is one)
- MBA\* is optimal (finds the shortest solution if heuristic is admissible)

## Solution 1: Multibody A\*

- Branching factor is  $b^k$  where  $b$  is the branching factor of one agent



## Solution 1: Multibody A\*

- Branching factor is  $b^k$  where  $b$  is the branching factor of one agent
- We say the solution depth is  $d$  if there is a (shortest) sequence of  $d$  joint actions that lead all agent to their respective goals


## Solution 1: Multibody A\*

- Branching factor is  $b^k$  where  $b$  is the branching factor of one agent
- We say the solution depth is  $d$  if there is a (shortest) sequence of  $d$  joint actions that lead all agent to their respective goals
- Worst case: We have to check all search tree branches of length  $d$ , i.e., expand  $b^{kd}$  states

## Solution 1: Multibody A\*


- Branching factor is  $b^k$  where  $b$  is the branching factor of one agent
- We say the solution depth is  $d$  if there is a (shortest) sequence of  $d$  joint actions that lead all agent to their respective goals
- Worst case: We have to check all search tree branches of length  $d$ , i.e., expand  $b^{kd}$  states
- Time/space complexity:  $O(b^{kd})$  (upper bound)

## Solution 1: Multibody A\*





- In a grid domain  $b = 5$ : Move(N), Move(E), Move(S), Move(W), Noop
- $k = 4$  agent
- Solution depth  $d = 7$
- Heuristic guides search perfectly!



## Solution 1: Multibody A\*





## Solution 1: Multibody A\*



## Solution 1: Multibody A\*





# Solution 1: Multibody A\*






## Solution 1: Multibody A\*




We expand  $d = 7$  states, but generate  $db^k = 7 \times 5^4 \approx 4300$  states(!)

If each agent could plan for itself, we would generate only  
 $dbk = 7 \times 5 \times 4 = 140$  states

## Solution 1: Multibody A\*



## Solution 1: Multibody A\*



Worst case: All search tree paths branches of length 3 are expanded before we find the solution of length 4 (one agent has to wait), i.e., all  $m^2 = b^k$  nodes are expanded.

## Solution 2: ???

- Multibody vs multi-agent planning: Simplicity vs efficiency (avoiding exponential blowup)

## Solution 2: ???

- Multibody vs multi-agent planning: Simplicity vs efficiency (avoiding exponential blowup)
- "Best" case (heuristic guides search perfectly):
  - Reduce branching factor
  - Solve  $k$  problems of time/space complexity  $O(db)$
  - ... instead of one problem of  $O(db^k)$  (140 vs 4300)

## Solution 2: ???

- Multibody vs multi-agent planning: Simplicity vs efficiency (avoiding exponential blowup)
- "Best" case (heuristic guides search perfectly):
  - Reduce branching factor
  - Solve  $k$  problems of time/space complexity  $O(db)$
  - ... instead of one problem of  $O(db^k)$  (140 vs 4300)
- Worst case:
  - Solve  $k$  problems of time/space complexity  $O(b^d) + \text{overhead}$
  - ... instead of one problem of  $O(b^{kd})$

## Solution 2: ???

- Multibody vs multi-agent planning: Simplicity vs efficiency (avoiding exponential blowup)
- "Best" case (heuristic guides search perfectly):
  - Reduce branching factor
  - Solve  $k$  problems of time/space complexity  $O(db)$
  - ... instead of one problem of  $O(db^k)$  (140 vs 4300)
- Worst case:
  - Solve  $k$  problems of time/space complexity  $O(b^d) + \text{overhead}$
  - ... instead of one problem of  $O(b^{kd})$
- Can we use CDPS (cooperative distributed problem solving) for MAPF?
  1. Problem decomposition
  2. Subproblem solution
  3. Solution synthesis

## Solution 2: ???

Discuss: How do we ensure the plans of each agent are compatible?




Illustration from August Møbius' thesis *Multi-Agent Pathfinding*



## Solution 2: Space-time A\*

Space-time A\* (STA\*): “A\* with time constraints”

### Definition

A constraint  $(a, v, t)$  precludes agent  $a$  from being in location  $v \in V$  at time  $t$




Illustration from August Møbius' thesis *Multi-Agent Pathfinding*

## Solution 2: Space-time A\*


```
def SpaceTimeAStar(problem, agents):  
    solution = dict()  
    constraints = []  
    k = len(agents)  
    for i in range(k):  
        plan = AStar(problem[agents[i]], constraints)  
        constraints += [(agents[j], v, t)  
                        for j in range(i+1, k)  
                        for t, v in enumerate(plan)]  
        solution[agent] = plan  
    return solution
```

Assume a sub-routine  $AStar(p, c)$  with an admissible heuristic that can solve single-agent problem instances given



- a single-agent problem instance,  $p$ , and
- a (possibly empty) list of *constraints*,  $c$ .

Here, the constraints  $c$  are the solutions of the previous agents




## Solution 2: Space-time A\*



## Solution 2: Space-time $A^*$



## Solution 2: Space-time A\*



## Solution 2: Space-time A\*

Questions:

- Is it optimal? Can you give an example where we're not guaranteed to find the shortest solution?
- Is it complete? Can you give an example where we're not guaranteed to find a solution when there is one?

## Solution 2: Space-time A\*

Questions:

- Is it optimal? Can you give an example where we're not guaranteed to find the shortest solution?
- Is it complete? Can you give an example where we're not guaranteed to find a solution when there is one?

We have serialisable subgoals, but STA\* doesn't choose subgoals intelligently





Illustration from August Møbius' thesis *Multi-Agent Pathfinding*

## Solution 3: ...

Can we get the best of both worlds?

- Agents plan independently if they can
- Agents plan as a multibody if they conflict



## Solution 3: ...

Can we get the best of both worlds?

- Agents plan independently if they can
- Agents plan as a multibody if they conflict

### Definition

$(a_i, a_j, v, t)$  is a conflict of agents  $a_i$  and  $a_j$  both planning to be in location  $v \in V$  at time  $t$ .

## Solution 3: ...

Can we get the best of both worlds?

- Agents plan independently if they can
- Agents plan as a multibody if they conflict

### Definition

$(a_i, a_j, v, t)$  is a conflict of agents  $a_i$  and  $a_j$  both planning to be in location  $v \in V$  at time  $t$ .

Algorithm called *Independence detection* (ID)

## Solution 3: Independence detection

```
def IndependenceDetection(problem, agents):  
    solution = dict()  
    groups = [[agent] for agent in agents]  
    for group in groups:  
        solution[group] = AStar(problem[group], [])  
    while not IsValid(solution):  
        # Merge conflicting groups  
        (group1, group2, v, t) = FindFirstConflict(solution)  
        del solution[group1]  
        del solution[group2]  
        merged = [group1, group2]  
        groups.append(merged)  
  
        # Plan for the merged group  
        solution[merged] = AStar(problem[merged], [])  
    return solution
```





$(group_1, group_2, v, t) = \text{FindFirstConflict}(\text{solution})$  is a conflict of (multibody) agents  $group_1$  and  $group_2$  because they both plan to be in location  $v$  at time  $t$ .

## Solution 3: Independence detection

- Optimal and complete
- In the worst case, we end up with a group containing all agents

## Solution 3: Independence detection

- Optimal and complete
- In the worst case, we end up with a group containing all agents



## Solution 4: Conflict based search

- Use reservations (like Space-time A\*)
  - but without blocking
- Re-plan in case of conflicts (like Independence Detection)
  - but without grouping

## Solution 4: Conflict based search


- Use reservations (like Space-time A\*)
  - but without blocking
- Re-plan in case of conflicts (like Independence Detection)
  - but without grouping
- Idea:
  - Use a high-level *constraint tree* to keep track of constraints
  - Use a low-level A\* search for each agent (taking constraints into account)

## Solution 4: Conflict based search


- Use reservations (like Space-time A\*)
  - but without blocking
- Re-plan in case of conflicts (like Independence Detection)
  - but without grouping
- Idea:
  - Use a high-level *constraint tree* to keep track of constraints
  - Use a low-level A\* search for each agent (taking constraints into account)
- Conflict based search (CBS)



## Solution 4: Conflict based search




## Solution 4: Conflict based search



$(\bullet, \bullet, D5, 4)$

Cost:  $8 + 7$


## Solution 4: Conflict based search



$(\bullet, D5, 4)$

Cost:  $9 + 7$

## Solution 4: Conflict based search




$(\bullet, D5, 4)$


Cost:  $8 + 8$

# Conflicts

Vertex conflict




Edge conflict




# Conflicts


Vertex conflict




Edge conflict



Follow conflict




## Solution 4: Conflict based search




Cost: 8 + 7

$(\bullet, D5, 4)$



Cost: 9 + 7

$(\bullet, D5, 4)$



Cost: 8 + 8

## Solution 4: Conflict based search

Cost functions:

$$\text{FLOWTIME}(\pi) = \sum_{i=1}^k |\pi_i|$$

$$\text{MAKESPAN}(\pi) = \max_{i=1}^k |\pi_i|$$



## Solution 4: Conflict based search


Cost functions:

$$\text{FLOWTIME}(\pi) = \sum_{i=1}^k |\pi_i|$$

$$\text{MAKESPAN}(\pi) = \max_{i=1}^k |\pi_i|$$


FUEL: Like FLOWTIME, but Noop is free

## Solution 4: Conflict based search



Cost:  $8 + 7$

$(\bullet, D5, 4)$




Cost:  $9 + 7$

$(\bullet, D5, 4)$



Cost:  $8 + 8$


## Solution 4: Conflict based search



Cost: 8 + 7


$(\bullet, D5, 4)$

$(\bullet, D5, 4)$



Cost: 9 + 7

$(\bullet, D5, 4)$



Cost: 8 + 8

$(\bullet, D5, 4)$



Cost: 9 + 8



Cost: 8 + 9

## Solution 4: Conflict based search

```
def ConflictBasedSearch(problem, agents, cost):
    root.constraints = []
    for a in agents:
        root.solution[a] = AStar(problem[a])
    root.cost = cost(root.solution)
    frontier = PriorityQueue(root, lambda n: n.cost)
    while not frontier.empty():
        node = frontier.get() # get lowest cost
        if IsValid(node.solution):
            return node.solution
        (a1, a2, v, t) = FindFirstConflict(node.solution)
        for a in [a1, a2]:
            m = n.copy()
            m.constraints.append((a, v, t))
            m.solution[a] = AStar(problem[a],
                                     m.constraints)
            m.cost = cost(m.solution)
            if m.cost < infinity: # solution is found
                frontier.put(m)
```

## Solution 4: Conflict based search

- CBS is complete and optimal because we're exploring all combinations of constraints to find a shortest solution for each agent

## Solution 4: Conflict based search


- CBS is complete and optimal because we're exploring all combinations of constraints to find a shortest solution for each agent
- Is CBS always better than MBA\*, ID etc.?

## Solution 4: Conflict based search

- CBS is complete and optimal because we're exploring all combinations of constraints to find a shortest solution for each agent
- Is CBS always better than MBA\*, ID etc.?
- No, not if two agents are *strongly coupled*, i.e., have many conflicts among them


## Solution 4: Conflict based search

- CBS is complete and optimal because we're exploring all combinations of constraints to find a shortest solution for each agent
- Is CBS always better than MBA\*, ID etc.?
- No, not if two agents are *strongly coupled*, i.e., have many conflicts among them
- In an open  $n \times n$  grid there are  $2^{n/2}$  optimal paths between opposite corners





## Solution 4: Conflict based search

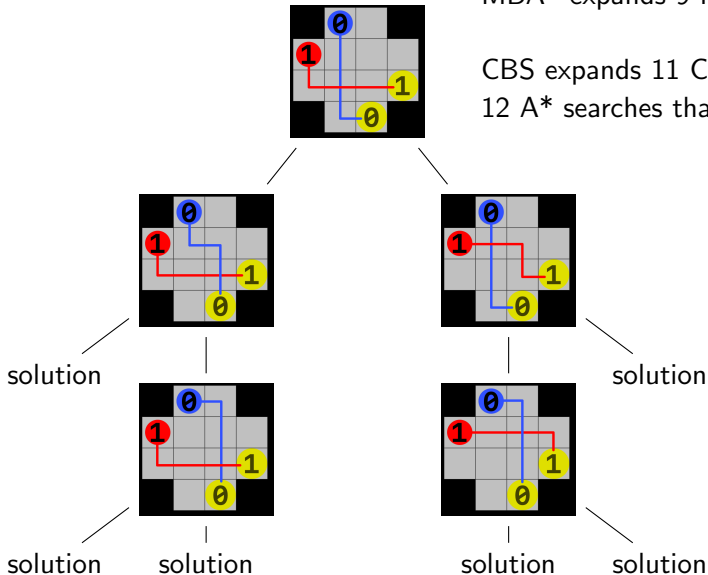


## Solution 4: Conflict based search

MBA\* expands 9 nodes

CBS expands 11 CT nodes:

12 A\* searches that expand 7 nodes



## Solution 4.5: Meta-agent CBS

- Improvement: Meta-agent CBS (MA-CBS)
- If there are more than  $x$  conflicts between  $a_i$  and  $a_j$ , merge them into one *meta-agent*
- Same idea as ID and the Martha project

## Solution 4.5: Meta-agent CBS

- Improvement: Meta-agent CBS (MA-CBS)
- If there are more than  $x$  conflicts between  $a_i$  and  $a_j$ , merge them into one *meta-agent*
- Same idea as ID and the Martha project
- We can tune the parameter  $x$ :
  - $\text{MA-CBS}(\infty) = \text{CBS}$
  - $\text{MA-CBS}(0) = \text{ID}$

## Solution 4.5: Meta-agent CBS

```
def ConflictBasedSearch(problem, agents, cost, x):
    conflicts = dict()
    root.constraints = []
    for a in agents:
        root.solution[a] = AStar(problem[a])
    root.cost = cost(root.solution)
    frontier = PriorityQueue(root, lambda n: n.cost)
    while not frontier.empty():
        node = frontier.get() # get lowest cost
        if IsValid(node.solution):
            return node.solution
        (a1, a2, v, t) = FindFirstConflict(node.solution)
        if conflicts[(a1,a2)] > x:
            m = merge(a1, a2, constraints)
            m.solution[a] = AStar(problem[(a1,a2)]
                                   m.constraints)
            m.cost = cost(m.solution)
            if m.cost < infinity: # solution is found
                frontier.put(m)
        else:
            conflicts[(a1,a2)] += 1 # then do as before
```

# Summary

	<b>Pros</b>	<b>Cons</b>
MBA*	It works!	But inefficient

# Summary

	<b>Pros</b>	<b>Cons</b>
MBA*	It works!	But inefficient
STA*	It's efficient!	But not complete, not optimal

# Summary

	<b>Pros</b>	<b>Cons</b>
MBA*	It works!	But inefficient
STA*	It's efficient!	But not complete, not optimal
ID	Complete and optimal!	But corridors degrade performance



# Summary

	<b>Pros</b>	<b>Cons</b>
MBA*	It works!	But inefficient
STA*	It's efficient!	But not complete, not optimal
ID	Complete and optimal!	But corridors degrade performance
CBS	Corridors, no problem!	But some agent merging again plz?

# Summary

	Pros	Cons
MBA*	It works!	But inefficient
STA*	It's efficient!	But not complete, not optimal
ID	Complete and optimal!	But corridors degrade performance
CBS	Corridors, no problem!	But some agent merging again plz?
MA-CBS	Have some merge!	But what's a good $x$ ?

## Now what?

- See (Silver, 2005) for STA\*
- See (Standley, 2012) for ID
- CBS was introduced in (Sharon et al., 2015) and compares it to other state-of-the-art algorithms
  - Benchmark against other algorithms
  - Focus on completeness and optimality

## Now what?

- See (Silver, 2005) for STA\*
- See (Standley, 2012) for ID
- CBS was introduced in (Sharon et al., 2015) and compares it to other state-of-the-art algorithms
  - Benchmark against other algorithms
  - Focus on completeness and optimality
- This won't solve the hospital domain, but hope you can take inspiration from the analysis and comparisons of algorithms



## References

- Guni Sharon, Roni Stern, Ariel Felner, and Nathan R. Sturtevant.  
Conflict-based search for optimal multi-agent pathfinding. *Artificial Intelligence*, 219:40–66, 2015. ISSN 0004-3702. doi:  
<https://doi.org/10.1016/j.artint.2014.11.006>. URL <https://www.sciencedirect.com/science/article/pii/S0004370214001386>.
- David Silver. Cooperative pathfinding. In *Proceedings of the First AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, page 117–122, 2005.
- Trevor Scott Standley. Independence detection for multi-agent pathfinding problems. In *Workshops at the Twenty-Sixth AAAI Conference on Artificial Intelligence*, 2012.